



## Chapter 19: Optimization

---

### CS-6360 Database Design

Chris Irwin Davis, Ph.D.

**Email:** [chrisirwindavis@utdallas.edu](mailto:chrisirwindavis@utdallas.edu)

**Phone:** (972) 883-3574

**Office:** ECSS 4.705

- **Query optimization:**
  - The process of choosing a suitable execution strategy for processing a query.
- Two high-level approaches:
  - **Heuristic Estimate**
  - **Cost-based Estimate**

## 19.1 – Query Trees and Heuristics for Query Optimization

- Process for heuristics optimization
  1. The parser of a high-level query generates an initial internal representation;
  2. Apply heuristics rules to optimize the internal representation.
  3. A query execution plan is generated to execute groups of operations based on the access paths available on the files involved in the query.
- The main heuristic is to first apply the operations that reduce the size of intermediate results.
  - e.g., Apply **SELECT** and **PROJECT** operations before applying the **JOIN** or other binary operations.

- Example:
  - For every project located in 'Stafford', retrieve the project number, the controlling department number and the department manager's last name, address and birthdate.
- Relation algebra:


$$\pi_{Pnumber, Dnum, Lname, Address, Bdate} (((\sigma_{Plocation='Stafford'}(PROJECT)) \bowtie_{Dnum=Dnumber} (DEPARTMENT)) \bowtie_{Mgr\_ssn=Ssn} (EMPLOYEE))$$

- SQL query:

**Q2: SELECT** P.Pnumber, P.Dnum, E.Lname, E.Address, E.Bdate  
**FROM** PROJECT **AS** P, DEPARTMENT **AS** D, EMPLOYEE **AS** E  
**WHERE** P.Dnum=D.Dnumber **AND** D.Mgr\_ssn=E.Ssn **AND**  
P.Plocation= 'Stafford';

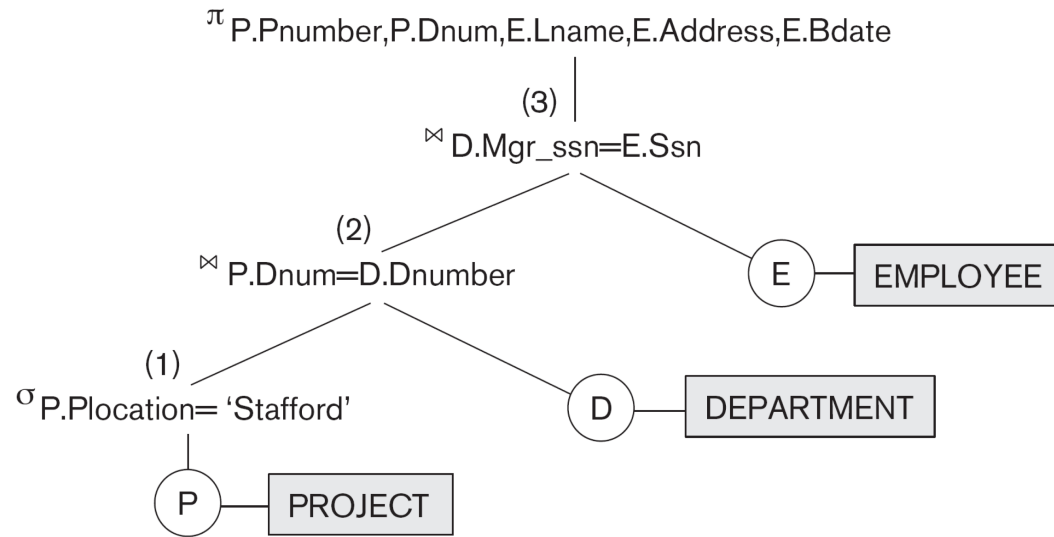
---

# Using Heuristics in Query Optimization




# Using Heuristics in Query Optimization

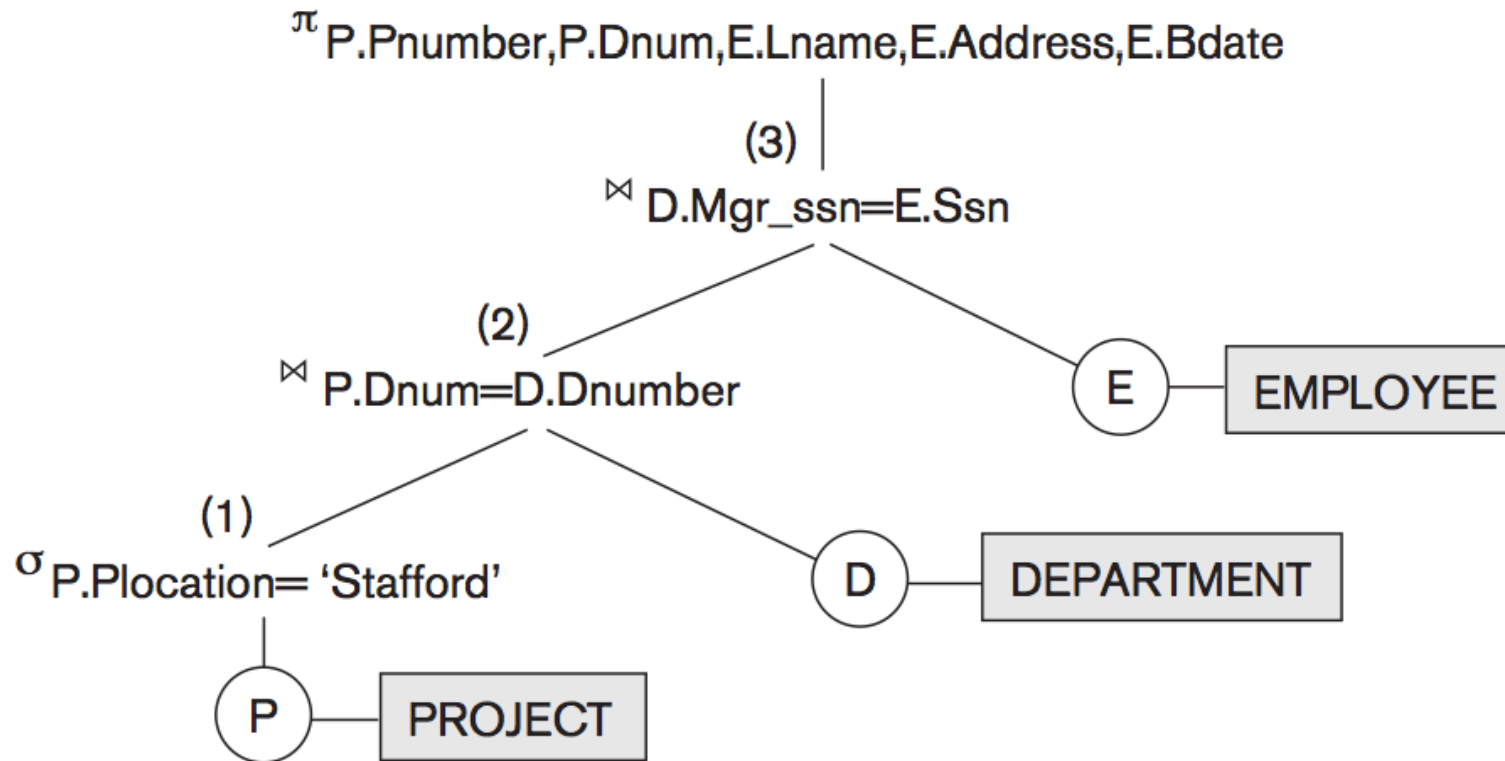
(a)



(b)



# Using Heuristics in Query Optimization



Query tree corresponding to the relational algebra expression for Q2.



# Using Heuristics in Query Optimization


---



**Q: SELECT** Lname  
**FROM** EMPLOYEE, WORKS\_ON, PROJECT  
**WHERE** Pname='Aquarius' **AND** Pnumber=Pno **AND** Essn=Ssn  
**AND** Bdate > '1957-12-31';

# Using Heuristics in Query Optimization

**Q: SELECT** Lname  
**FROM** EMPLOYEE, WORKS\_ON, PROJECT  
**WHERE** Pname='Aquarius' **AND** Pnumber=Pno **AND** Essn=Ssn  
**AND** Bdate > '1957-12-31';



Initial (canonical) query tree for SQL query Q2.


- Heuristic Optimization of Query Trees:
  - The same query could correspond to many different relational algebra expressions — and hence many different query trees.
  - The task of heuristic optimization of query trees is to find a **final query tree** that is efficient to execute.
- Example:

```
Q:  SELECT  Lname
      FROM    EMPLOYEE, WORKS_ON, PROJECT
      WHERE   Pname='Aquarius' AND Pnumber=Pno AND Essn=Ssn
            AND Bdate > '1957-12-31';
```

# Left-deep Tree Examples

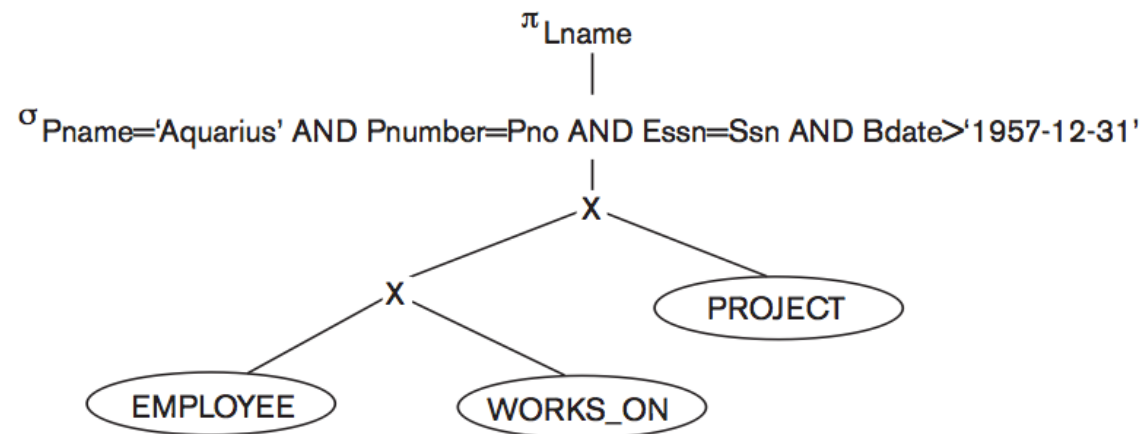
**Figure 19.7**

Two left-deep (JOIN) query trees.



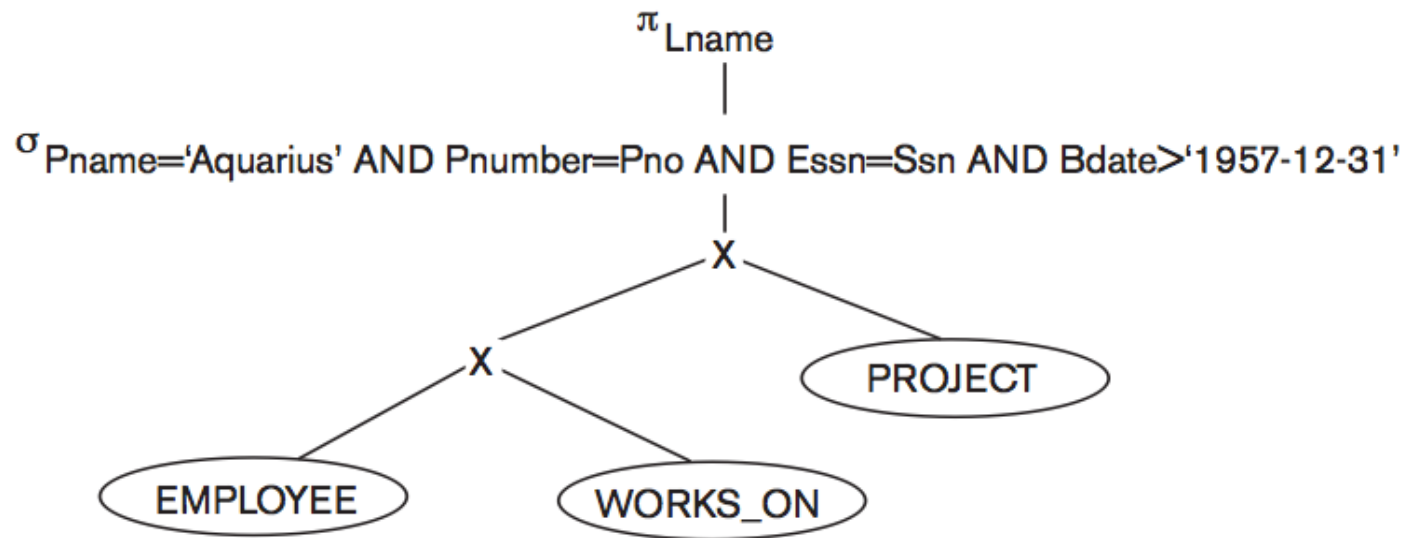
**Figure 19.5 (a)**

Query Q



# Steps in converting a query tree during heuristic optimization

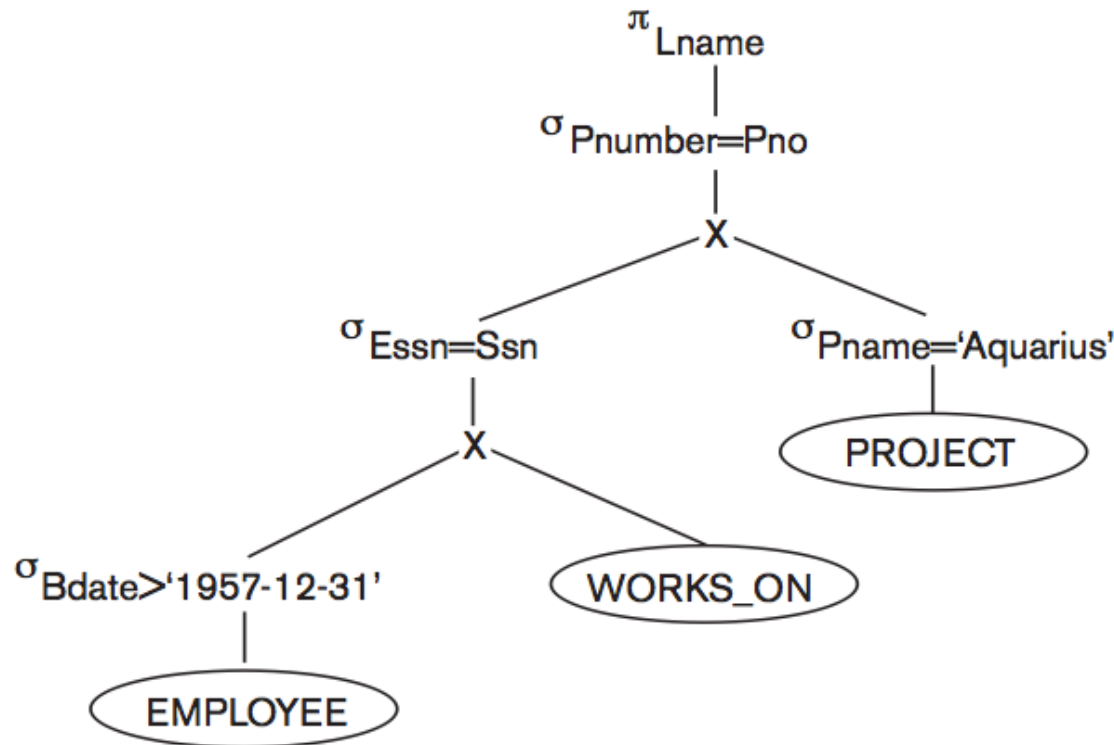
(a)



(a) Initial (canonical) query tree for SQL query Q.

# Steps in converting a query tree during heuristic optimization


(b)



(b) Moving SELECT operations down the query tree.


# Steps in converting a query tree during heuristic optimization

(c)



(c) Applying the more restrictive **SELECT** operation first.

# Steps in converting a query tree during heuristic optimization




(d) Replacing CARTESIAN PRODUCT + SELECT with JOIN operations.



# Steps in converting a query tree during heuristic optimization

(e)



(e) Moving **PROJECT** operations down the query tree.

## 19.8 – Using Selectivity and Cost Estimates in Query Optimization

# Using Selectivity and Cost Estimates in Query Optimization

---



- **Cost-based query optimization:**
  - Estimate and compare the costs of executing a query using different execution strategies and choose the strategy with the lowest cost estimate.
  - (Compare to heuristic query optimization)
- **Considerations and Issues**
  - Cost function
  - Number of execution strategies to be considered

# Using Selectivity and Cost Estimates in Query Optimization

---



- **Cost Components for Query Execution**
  1. Access cost to secondary storage
  2. Storage cost
  3. Computation cost
  4. Memory usage cost
  5. Communication cost
  
- Note: Different database systems may focus on different cost components.

- 1) Access cost to secondary storage
  - This is the cost of transferring (reading and writing) data blocks between secondary disk storage and main memory buffers.
  - This is also known as disk I/O (input/output) cost.
  - The cost of searching for records in a disk file depends on the type of access structures on that file, such as ordering, hashing, and primary or secondary indexes.
  - In addition, factors such as whether the file blocks are allocated contiguously on the same disk cylinder or scattered on the disk affect the access cost.

- 2) Disk storage cost
  - This is the cost of storing any intermediate files on disk that are generated by an execution strategy for the query.
- 3) Computation cost
  - This is the cost of performing in-memory operations on the records within the data buffers during query execution.
  - Such operations include searching for and sorting records, merging records for a join or a sort operation, and performing computations on field values.
  - This is also known as *CPU (central processing unit)* cost.

- **4) Memory usage cost**
  - This is the cost pertaining to the number of main memory buffers needed during query execution.
- **5) Communication cost**
  - This is the cost of shipping the query and its results from the database site to the site or terminal where the query originated.
  - In distributed databases (see Chapter 25), it would also include the cost of transferring tables and results among various computers during query evaluation.

# Using Selectivity and Cost Estimates in Query Optimization



- Catalog Information used in Cost Functions
  - To estimate the costs of various execution strategies, we must keep track of any information that is needed for the cost functions.
  - This information may be stored in the DBMS catalog, where it is accessed by the query optimizer.
  - **First**, we must know the size of each file. For a file whose records are all of the same type, the
    - **number of records (tuples) ( $r$ )**,
    - the (average) **record size ( $R$ )**, and
    - the **number of file blocks ( $b$ )** (or close estimates of them) are needed.
    - The **blocking factor ( $bfr$ )** for the file may also be needed.



# Using Selectivity and Cost Estimates in Query Optimization

---



- Catalog Information used in Cost Functions
- Information about the size of a **file**
  - $B$  – block size (in kb)
  - $R$  – record size (in kb)
  - $r$  – number of records (tuples)
  - $b$  – number of blocks
  - $bfr$  – blocking factor ( $\lfloor B / R \rfloor$  records per block)

# Using Selectivity and Cost Estimates in Query Optimization



- Catalog Information used in Cost Functions
- Information about **indices** and **indexing attributes** of a file
  - $r$  – Number of tuples (rows)
  - $sl$  – Selectivity of an attribute
    - fraction of records satisfying an equality condition
  - $s$  – Selection cardinality of an attribute. ( $s = sl * r$ )
  - $d$  – Number of distinct values of an attribute
  - $x$  – Number of levels of each multilevel index
  - $b_{I1}$  – Number of first-level index blocks
    - i.e. “root” level

# Using Selectivity and Cost Estimates in Query Optimization

- Catalog Information used in Cost Functions
- Information about **joins**
  - $|R \bowtie S|$ ,  $|R * S|$ ,  $|R \times S|$  – join cardinality
  - $js$  – join selectivity
    - The ratio of the expected size of the join result divided by the maximum size  $n_R * n_S$  of the join

## Examples of Cost Functions for **SELECT**

- **S1.** Linear search (brute force) approach
  - $C_{S1a} = b$ ;
  - For an equality condition on a key,  $C_{S1a} = (b/2)$  if the record is found; otherwise  $C_{S1a} = b$ .
- **S2.** Binary search:
  - $C_{S2} = \log_2 b + \lceil (s/bfr) \rceil - 1$
  - For an equality condition on a unique (key) attribute,  $C_{S2} = \log_2 b$
- **S3.** Using a primary index (S3a) or hash key (S3b) to retrieve a single record
  - $C_{S3a} = x + 1$ ;  $C_{S3b} = 1$  for static or linear hashing;
  - $C_{S3b} = 1$  for extendible hashing;

## Examples of Cost Functions for SELECT (contd.)

- **S4.** Using an ordering index to retrieve multiple records:
  - For the comparison condition on a key field with an ordering index,  $C_{S4} = x + (b/2)$
- **S5.** Using a clustering index to retrieve multiple records:
  - $C_{S5} = x + \lceil (s/bfr) \rceil$
- **S6.** Using a secondary (B+-tree) index:
  - For an equality comparison,  $C_{S6a} = x + s$ ;
  - For an comparison condition such as  $>$ ,  $<$ ,  $>=$ , or  $<=$ ,
  - $C_{S6a} = x + (b_{I1}/2) + (r/2)$

## Examples of Cost Functions for SELECT (contd.)

- **S7. Conjunctive selection:**
  - Use either S1 or one of the methods S2 to S6 to solve.
  - For the latter case, use one condition to retrieve the records and then check in the memory buffer whether each retrieved record satisfies the remaining conditions in the conjunction.
- **S8. Conjunctive selection using a composite index:**
  - Same as S3a, S5 or S6a, depending on the type of index.
- Examples of using the cost functions.

## Examples of Cost Functions for JOIN (contd.)

- **J1. Nested-loop join:**
  - Using  $R$  for outer loop
  - $C_{J1} = b_R + (b_R * b_S) + ((js * |R| * |S|) / bfr_{RS})$
- **J2. Single-loop join (using an access structure to retrieve the matching record(s))**
  - If an index exists for the join attribute  $B$  of  $S$  with index levels  $x_B$ , we can retrieve each record  $s$  in  $R$  and then use the index to retrieve all the matching records  $t$  from  $S$  that satisfy  $t[B] = s[A]$ .
  - The cost depends on the type of index.

# Using Selectivity and Cost Estimates in Query Optimization

## Examples of Cost Functions for JOIN (contd.)

### ■ J2. Single-loop join (contd.)

#### ■ For a **secondary index** on $S$ ,

$$\square C_{J2a} = b_R + (|R| * (x_B + 1 + s_B)) + ((j_s * |R| * |S|) / bfr_{RS})$$

#### ■ For a **clustering index** on $S$ ,

$$\square C_{J2b} = b_R + (|R| * (x_B + (s_B / bfr_B))) + ((j_s * |R| * |S|) / bfr_{RS});$$

#### ■ For a **primary index** on $S$ ,

$$\square C_{J2c} = b_R + (|R| * (x_B + 1)) + ((j_s * |R| * |S|) / bfr_{RS})$$

#### ■ If a hash key exists for one of the two join attributes — $B$ of $S$

$$\square C_{J2d} = b_R + (|R| * h) + ((j_s * |R| * |S|) / bfr_{RS})$$

### ■ J3. Sort-merge join:

$$\square C_{J3a} = C_S + b_R + b_S + ((j_s * |R| * |S|) / bfr_{RS})$$

$$\square (C_S: \text{Cost for sorting files})$$



## Example: Join Cost

- Suppose that we have the EMPLOYEE file described in the example in the previous section (of the textbook, per p.717), and
- Assume that the DEPARTMENT file in Figure 3.5 consists of  $r_D = 125$  records stored in  $b_D = 13$  disk blocks
  - EMPLOYEE cost values are from the table in Figure 19.8 (book never says so)
  - DEPARTMENT cost values are given (different than Figure 19.8)
- Consider the following two join operations:
  - OP6:  $\text{EMPLOYEE} \bowtie_{\text{Dno=Dnumber}} \text{DEPARTMENT}$
  - OP7:  $\text{DEPARTMENT} \bowtie_{\text{Mgr\_ssn=Ssn}} \text{EMPLOYEE}$

## Example: Join Cost

- Suppose that we
  - primary index on Dnumber of DEPARTMENT with  $x_{\text{Dnumber}} = 1$  level
  - secondary index on Mgr\_ssn of DEPARTMENT with selection cardinality  $s_{\text{Mgr\_ssn}} = 1$  and levels  $x_{\text{Mgr\_ssn}} = 2$ .
- Assume that the join selectivity for OP6 is
$$j_{\text{SOP6}} = (1 / |\text{DEPARTMENT}|) = 1 / 125$$
  - because Dnumber is a key of DEPARTMENT.
- Also assume that the blocking factor for the resulting join file is
$$bfr_{\text{ED}} = 4 \text{ records per block}$$

1. Using method J1 with EMPLOYEE as outer loop:

$$\begin{aligned} C_{J1} &= b_E + (b_E * b_D) + ((js_{OP6} * r_E * r_D)/bfr_{ED}) \\ &= 2000 + (2000 * 13) + (((1/125) * 10,000 * 125)/4) = 30,500 \end{aligned}$$

2. Using method J1 with DEPARTMENT as outer loop:

$$\begin{aligned} C_{J1} &= b_D + (b_E * b_D) + ((js_{OP6} * r_E * r_D)/bfr_{ED}) \\ &= 13 + (13 * 2000) + (((1/125) * 10,000 * 125)/4) = 28,513 \end{aligned}$$

3. Using method J2 with EMPLOYEE as outer loop:

$$\begin{aligned} C_{J2c} &= b_E + (r_E * (x_{Dnumber} + 1)) + ((js_{OP6} * r_E * r_D)/bfr_{ED}) \\ &= 2000 + (10,000 * 2) + (((1/125) * 10,000 * 125)/4) = 24,500 \end{aligned}$$

4. Using method J2 with DEPARTMENT as outer loop:

$$\begin{aligned} C_{J2a} &= b_D + (r_D * (x_{Dno} + s_{Dno})) + ((js_{OP6} * r_E * r_D)/bfr_{ED}) \\ &= 13 + (125 * (2 + 80)) + (((1/125) * 10,000 * 125)/4) = 12,763 \end{aligned}$$